Nguyen Thanh Duy Tan

AI Engineer Fresher

Tan Binh District, Ho Chi Minh City | 0339900276 | tanntd.2005@gmail.com | github.com/TanNguyen234 | linkedin.com/in/tannguyen234 | nguyentan234.vercel.app

Career Objective

Short-term: Apply AI/ML knowledge to real projects and improve practical engineering skills.
Long-term: Become an AI Engineer capable of building reliable ML and LLM-based systems.

Education

Ho Chi Minh City University of Natural Resources and Environment — Oct 2023 – Present
Major: Information Technology — GPA: 3.1/4.0

Skills

- **Machine Learning:** Scikit-learn, Pandas, NumPy, Imbalanced Data Handling, EDA.
- **Generative AI & NLP:** LLMs, LangGraph, LangChain, Multi-Agent Systems, RAG.
- **MLOps & Software Engineering:** FastAPI, Docker, CI/CD (GitHub Actions), Pytest, RESTful APIs.
- **Database & Storage:** MongoDB (Motor Async).
- **Tools:** Git, GitHub, Jupyter Notebook, Linux.
- **Soft Skills:** Problem-solving, system design thinking, self-researching technical docs.

Projects

CV AI Evaluation System (Generative AI & LLM) — 03/2026 – Present

- Full-stack, multi-agent AI system for automated resume parsing, profiling, and professional evaluation.
- Designed a LangGraph workflow to run Foundational, Professional, and Project evaluation phases concurrently.
- Integrated Tavily API (RAG) for market context and a Meta-Evaluator to check cross-agent logic consistency.
- Built an async FastAPI backend with SSE progress streaming and Base64 CV storage in MongoDB.
- Added a 5-requests/day rate limiter and AI Spam Validation to filter invalid files and reduce LLM API costs.

Tech Stack: Python, FastAPI, LangGraph, LangChain, LLMs (Gemini, Qwen), SSE, MongoDB, PyMuPDF4LLM.
GitHub: github.com/TanNguyen234/CV-Review-System
Demo: foxy0505-ai-cv-reviewer.hf.space/app

End-to-End Fraud Detection Pipeline (Traditional ML & MLOps) — 05/2026 – Present

- Production-ready machine learning service to detect credit card fraud with high accuracy.
- Conducted EDA and trained an XGBoost classifier for a highly imbalanced dataset with only 0.17% fraud cases.
- Packaged the model into a FastAPI app for single-transaction prediction and batch inference up to 100 transactions.
- Containerized the app with Docker and built a GitHub Actions CI/CD pipeline for Ruff linting, testing, and image deployment.
- Wrote unit and integration tests with Pytest, using Pydantic for strict API schema validation.

Tech Stack: Python, XGBoost, Pandas, Scikit-learn, FastAPI, Docker, GitHub Actions, Pytest.
GitHub: github.com/TanNguyen234/End-to-End-Fraud-Detection-Pipeline