

Nguyen Thanh Duy Tan

AI Engineer Fresher

Tan Binh District, Ho Chi Minh City | +84 339 900 276 | tanntd.2005@gmail.com | github.com/TanNguyen234 |
linkedin.com/in/tannguyen234 | nguyentan234.vercel.app

Career Objective

Short-term: Apply AI/ML knowledge to real projects and improve practical engineering skills.
Long-term: Become an AI Engineer capable of building reliable ML and LLM-based systems.

Education

Ho Chi Minh City University of Natural Resources and Environment Oct 2023 – Present
Major: Information Technology GPA: 3.0/4.0

Achievements

Academic Encouragement Scholarship – Good Classification, Academic Year 2024–2025.

Skills

- **Machine Learning:** Scikit-learn, Pandas, NumPy, Imbalanced Data Handling, EDA.
- **Generative AI & NLP:** LLMs, LangGraph, LangChain, Multi-Agent Systems, RAG.
- **MLOps & Software Engineering:** FastAPI, Docker, CI/CD (GitHub Actions), RESTful APIs.
- **Database & Storage:** MongoDB (Motor Async).
- **Tools:** Git, GitHub, Jupyter Notebook.
- **Soft Skills:** Problem-solving, system design thinking, self-researching technical docs.

Projects

VietLex Legal RAG System 07/2026 – Present

- Designed a five-stage legal retrieval flow: query rewriting, Qdrant hybrid search with BGE-M3 dense/BM25 sparse retrieval, RRF Top 15 fusion, and Cohere Top 3 reranking.
- Implemented a Qdrant semantic cache at a ≥ 0.96 similarity threshold, reducing repeated-query latency by 85% and avoiding unnecessary LLM calls.
- Integrated cleaned Vietnamese legal-document ingestion, NeMo Guardrails, SlowAPI limits (5 requests/minute/IP), Logfire tracing, and Ragas background evaluation.

Tech Stack: Python, FastAPI, LangChain, Qdrant, BGE-M3, Cohere Rerank, NeMo Guardrails, Ragas, Docker.
GitHub: github.com/TanNguyen234/VietLex-Tech-Spec

CV AI Evaluation System (Generative AI & LLM) 03/2026 – 05/2026

- Built an end-to-end pipeline for PDF extraction, candidate profiling, parallel foundational/professional/project evaluation, validation, and structured report generation.
- Orchestrated specialized agents with LangGraph, enriched scoring rubrics using Tavily market context, and used a Meta-Evaluator for cross-agent consistency checks.
- Delivered an async FastAPI service with SSE progress streaming, Motor-based MongoDB history, AI spam validation, and a 5-requests/day rate limit.

Tech Stack: Python, FastAPI, LangGraph, LangChain, LLMs (Gemini, Qwen), SSE, MongoDB, PyMuPDF4LLM.
GitHub: github.com/TanNguyen234/CV-Review-System | **Demo:** foxy0505-ai-cv-reviewer.hf.space/app

End-to-End Fraud Detection Pipeline (Traditional ML & MLOps) 05/2026 – 06/2026

- Performed EDA and built an imbalance-aware XGBoost training pipeline on transaction data containing only 0.17% fraud cases.
- Exposed the model through FastAPI with strict Pydantic schemas, supporting single predictions and batch inference of up to 100 transactions.
- Dockerized the service and configured GitHub Actions for Ruff linting, automated tests, image builds, and reproducible deployment.

Tech Stack: Python, XGBoost, Pandas, Scikit-learn, FastAPI, Docker, GitHub Actions.
GitHub: github.com/TanNguyen234/End-to-End-Fraud-Detection-Pipeline